

Statistisches Data Mining (StDM)

Woche 10

Oliver Dürr

Institut für Datenanalyse und Prozessdesign

Zürcher Hochschule für Angewandte Wissenschaften

oliver.duerr@zhaw.ch

Winterthur, 22 November 2016

No laptops, no phones, no problems




Multitasking senkt Lerneffizienz:

- **Keine Laptops im Theorie-Unterricht Deckel zu oder fast zu (Sleep modus)**

Overview of classification (until the end to the semester)

Classifiers



K-Nearest-Neighbors (KNN)

Logistic Regression

Linear discriminant analysis


Support Vector Machine (SVM)

Classification Trees

Neural networks NN

Deep Neural Networks (e.g. CNN, RNN)

...



Combining classifiers

Bagging

Boosting

Random Forest

Evaluation



Cross validation


Performance measures

ROC Analysis / Lift Charts

Theoretical Guidance / General Ideas

Bayes Classifier


Bias Variance Trade
off (Overfitting)



Feature Engineering

Feature Extraction

Feature Selection



SVM


Chapter 9 in ILSR

Note on notation in ISLR


- In ISLR they make an unusual distinction between Support Vector Classifier and Support Vector Machine (SVM).
- Here we call everything a SVM
 - Linear Separable Case
 - SVM with Penalty allowing misclassifications
 - SVM with Kernels

Support Vector Machine (SVM) - Basics

- Each observation \Leftrightarrow vector of values (p-Dimensional)
- SVM constructs a hyperplane to separate class members.



Welche Ebene?




Zeichnen Sie eine Linie, die die beiden Klassen möglichst gut trennt.

Begründen Sie Ihre Wahl

Support Vector Machine - Hyperplanes


- Each column vector can be viewed as a point in an p -dimensional space (p = number of features).
- A linear binary classifier constructs a hyperplane separating class members from non-members in this space.



...which one?


Support Vector Machine - Maximum Margin Hyperplane

- SVM choose a specific hyperplane among the many that can separate the data, namely the *maximum margin hyperplane*, which maximizes the distance from the hyperplane to the closest training point.
- The maximum margin hyperplane can be represented as a linear combination of (some) training points.



SVM - Support Vectors

- Training examples that lie far away from the hyperplane do not participate in its specification.
- Training examples that lie closest to the decision boundary between the two classes determine the hyperplane.
- These training examples are called the **support vectors**, since removing them would change the location of the separating hyperplane. They determine the classifier.



Mathematical Definition and Optimization (just sketch)

Formal:: Definition of a hyperplane

We assume that classes are separable

$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p = 0 \quad \text{Definition of Hyperplane}$$

Separating hyperplane for classes coded as $y=\pm 1$

$$\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip} > 0 \text{ if } y_i = 1,$$

$$\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip} < 0 \text{ if } y_i = -1.$$


Combining

$$y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) > 0$$

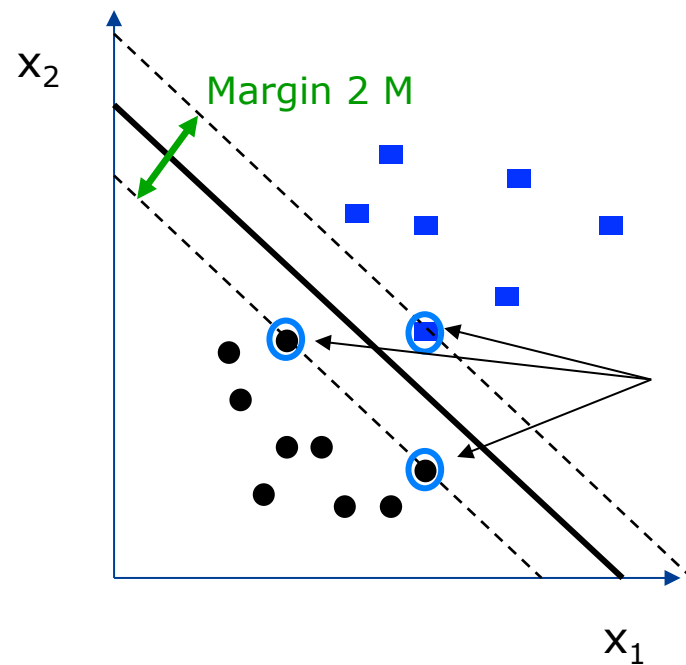
Note that this is only up to a constant (multiplication does not change anything)

→ Fix beta for components 1 to p:

$$\sum_{j=1}^p \beta_j^2 = 1, \quad \beta \text{ (for } j=1, \dots, p) \text{ is a normal vector}$$



Formal:: Definition of optimization problem



Support Vectors, have distance M to hyperplane

Intuitive Optimization

maximize M
 $\beta_0, \beta_1, \dots, \beta_p$

subject to $\sum_{j=1}^p \beta_j^2 = 1,$

$y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M \quad \forall i = 1, \dots, n.$

All vectors have at least distance M.
Support Vectors have = M.

Formal:: Reformulating the optimization problem

$$\begin{array}{ll}\text{maximize } M \\ \beta_0, \beta_1, \dots, \beta_p\end{array}$$

Intuitive Optimization

$$\text{subject to } \sum_{j=1}^p \beta_j^2 = 1,$$

$$y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M \quad \forall i = 1, \dots, n.$$

Can be reformulated using Lagrange multipliers to

$$L_D = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{k=1}^N \alpha_i \alpha_k y_i y_k x_i^T x_k$$

Technical Optimization

$$\text{subject to } \alpha_i \geq 0 \text{ and } \sum_{i=1}^N \alpha_i y_i = 0.$$

once we have calculated the α 's we can calculate β (1,...,p) via


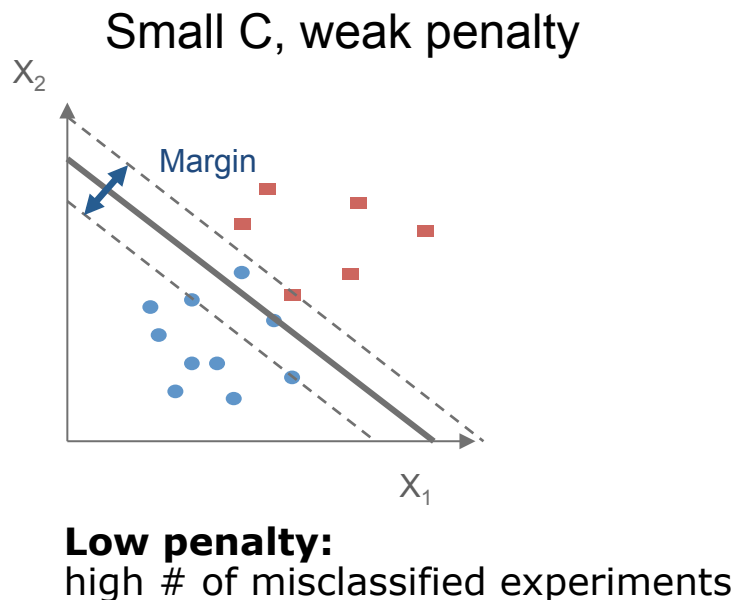
$$\beta = \sum_{i=1}^N \alpha_i y_i x_i$$

Only the inner product between the vectors of observations enters.


Opens the door to the kernel trick (see below)

SVM – Penalty (general idea)

- SVM may not be able to find any separating hyperplane at all, because the data contains untypical or mislabelled experiments.
- The problem can be addressed by using a *soft margin* that accepts some misclassifications of the training examples. The number of misclassifications is triggered by a *penalty factor* C .
- Sometimes a larger margin is worth having some misclassified observations



Formal:: SVM - Penalty



Introduction of slack variables ξ_i , for all observations (measured in units of M).

$$\text{maximize}_{\beta_0, \beta_1, \dots, \beta_p, \epsilon_1, \dots, \epsilon_n} M$$

Intuitive Optimization

$$\text{subject to } \sum_{j=1}^p \beta_j^2 = 1,$$

$$y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M(1 - \epsilon_i)$$

$$\epsilon_i \geq 0, \quad \sum_{i=1}^n \epsilon_i \leq C^*,$$

Finally this leads to the following equivalent optimization of L_D with constraints on α .

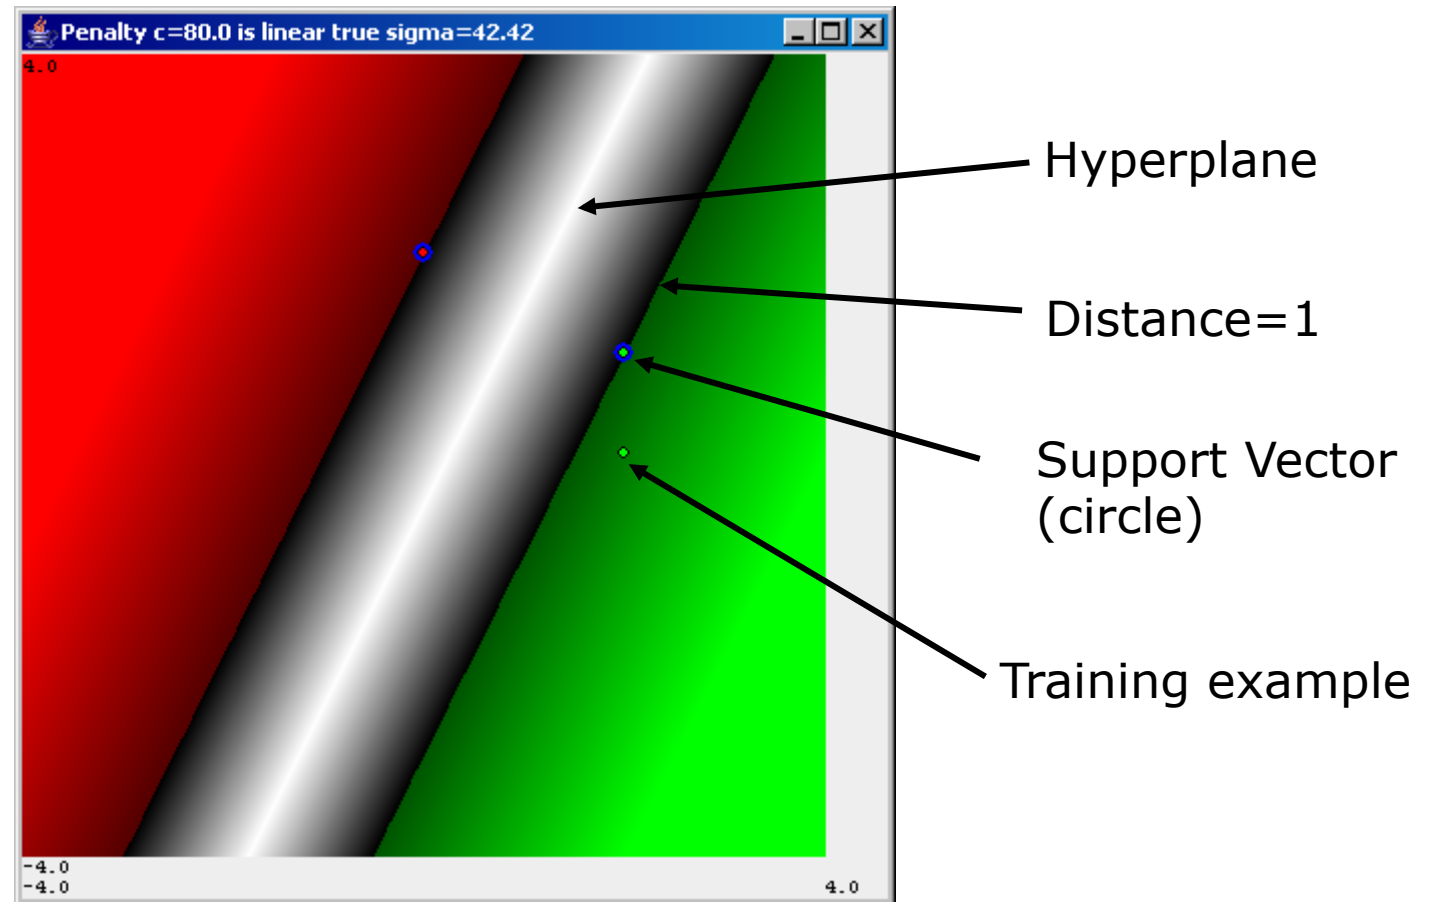
Technical Optimization ("dual form")

$$L_D = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{i'=1}^N \alpha_i \alpha_{i'} y_i y_{i'} x_i^T x_{i'} \quad 0 \leq \alpha_i \leq C \quad \sum_{i=1}^N \alpha_i y_i = 0$$


From α , β is obtained

Visualization of the parameter influence

Linear case effect of C



SVM – From low and high penalty



Very low c , nearly no penalty for misclassifications. Big margin. Let's increase c and see what happens.


Increase of c leads to convergence to a stable solution.

Why do we want a large margin?

- The margin controls the bias and variance
- Small margin (large C)
 - We expect that the margin depends more on the details of the concrete realization of the data. Hence: **large variance, small bias**
- Large margin (small C)
 - The margin depends less on the details of the concrete realization. Hence **small variance, large bias**

“Experimental” Observations (SVM) for gene expression

- Geneexpression: $p \gg N$
- C too low nearly no penalty for misclassification:
 - Overgeneralization (“don’t care”)
- C larger :
 - Converting to a stable solution.



Typical curve for gene expression
Misclassification rate as a function
of Log(C)

In general C is a hyper-parameter
which can be optimized (beware of
overfitting)

SVM in R (two classes)

```
library(e1071)
iris1 = iris[51:150,]
table(iris1$Species)
fit = svm(Species ~ ., data=iris1, kernel="linear", cost=10)
res = predict(fit, iris1)
sum(res == iris1$Species)
```

```
res_tune = tune(svm, Species ~ ., data=iris1,
kernel="linear", ranges = list(cost = c(0.1,1,10)))
summary(res_tune)
```

...


- Detailed performance results:

	cost	error	dispersion
1	0.1	0.04	0.05621827
2	1.0	0.04	0.03442652
3	10.0	0.04	0.03442652

Kernels

SVM - Non-separable data in the input space


- Some problems involve non-separable data for which there does not exist a hyperplane.
- The solution is to map the data into a higher-dimensional space and define a separating hyperplane there.





Note that this is often not the typical case.

Variable Transformation, make non-separable case separable

- Only a single variable x .
- Not separable by a point (hyperplane in 1D)



Take single variable x and x^2
Separable by a line (hyperplane in 2D)



View again in 1D

SVM - Feature space

- This higher-dimensional space is called the ***feature space*** as opposed to the input space.
- With an appropriately chosen feature space of sufficient dimensionality any consistent training set can be made separable.
- Example (last slide) $x \rightarrow (x, x^2)$
- In the program, one has to calculate

Optimization:

$$L_D = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{i'=1}^N \alpha_i \alpha_{i'} y_i y_{i'} x_i^T x_{i'} \quad 0 \leq \alpha_i \leq C \quad \sum_{i=1}^N \alpha_i y_i = 0$$



The only place where x enters

Kernel Trick

Optimization:

$$L_D = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{i'=1}^N \alpha_i \alpha_{i'} y_i y_{i'} x_i^T x_{i'} \quad 0 \leq \alpha_i \leq C \quad \sum_{i=1}^N \alpha_i y_i = 0$$

The only place where x enters

A handwritten derivation on a grid background showing the expansion of the dot product $x_i^T x_{i'}$. It is written as $(x_{i1}, x_{i2}, \dots, x_{ip}) \begin{pmatrix} x_{i'1} \\ x_{i'2} \\ \vdots \\ x_{i'p} \end{pmatrix} = \sum_{j=1}^p x_{ij} x_{i'j} =: K(x_i, x_{i'})$.

$$x_i^T x_{i'} = (x_{i1}, x_{i2}, \dots, x_{ip}) \begin{pmatrix} x_{i'1} \\ x_{i'2} \\ \vdots \\ x_{i'p} \end{pmatrix} = \sum_{j=1}^p x_{ij} x_{i'j} =: K(x_i, x_{i'})$$

- Kernel Trick:**


Replace: $K(x_i, x_{i'}) = \sum_{j=1}^p x_{ij} x_{i'j}$

With: $K(x_i, x_{i'}) = \sum_{j=1}^p x_{ij} x_{i'j} + \sum_{j=1}^p x_{ij}^2 x_{i'j}^2$

Is the same as explicitly making new features.

“Computed on the fly”

Hot topic in 1990's and early 2000s and still used



Kernel functions

- Instead of calculating the inner product, we calculate the kernel.
The following Kernels are commonly used:

$$K(x_i, x_{i'}) = \sum_{j=1}^p x_{ij} x_{i'j}$$

Identity (just the inner product)
In R 'linear kernel'

$$K(x_i, x_{i'}) = (1 + \sum_{j=1}^p x_{ij} x_{i'j})^d$$


Polynomial of degree d

$$K(x_i, x_{i'}) = \exp(-\gamma \sum_{j=1}^p (x_{ij} - x_{i'j})^2)$$


Gaussian, aka radial basis RBF.
Sometime $\gamma = 1/\sigma^2$

Kernels can also be used when data is not in the vector format. E.g. string kernels on text.

Example non-separable




Polynomial



RBF

SVM - Gaussian


- In a space in which the members of a class form one or more clusters, an accurate classifier might place a Gaussian around each cluster, thereby separating the clusters from the remaining space of non-class members.
- This effect can be accomplished by placing a Gaussian with a width (sigma) over each support vector in the training set.



$$K(\mathbf{X}, \mathbf{Y}) = \exp\left(\frac{-\|\tilde{\mathbf{X}} - \tilde{\mathbf{Y}}\|^2}{2\sigma^2}\right)$$


Visualization of the parameter influence

Gaussian Kernel effect of sigma



Gaussian Kernel in R

```
#####  
# Non-Linear Decision Boundary  
set.seed(1)  
x=matrix(rnorm(200*2), ncol=2)  
x[1:100,]=x[1:100,]+2  
x[101:150,]=x[101:150,]-2  
y=c(rep(1,150),rep(2,50))  
dat=data.frame(x=x,y=as.factor(y))  
  
require(manipulate)  
manipulate({  
  svmfit=svm(y ~ .,data=dat, kernel="radial",  
gamma=gamma, cost = cost)  
  plot(svmfit , dat) #Plotting  
}, gamma = slider(0.1,10), cost=slider(0.1,10))
```



Separation and dimensionality

Consider examples of 2 classes

Draw 2 points on a line. Can you always separate them?

Draw 3 points in a plane (not in a line!). Can you always separate them?

Imaging 4 points 3D, can you always separate them?

...

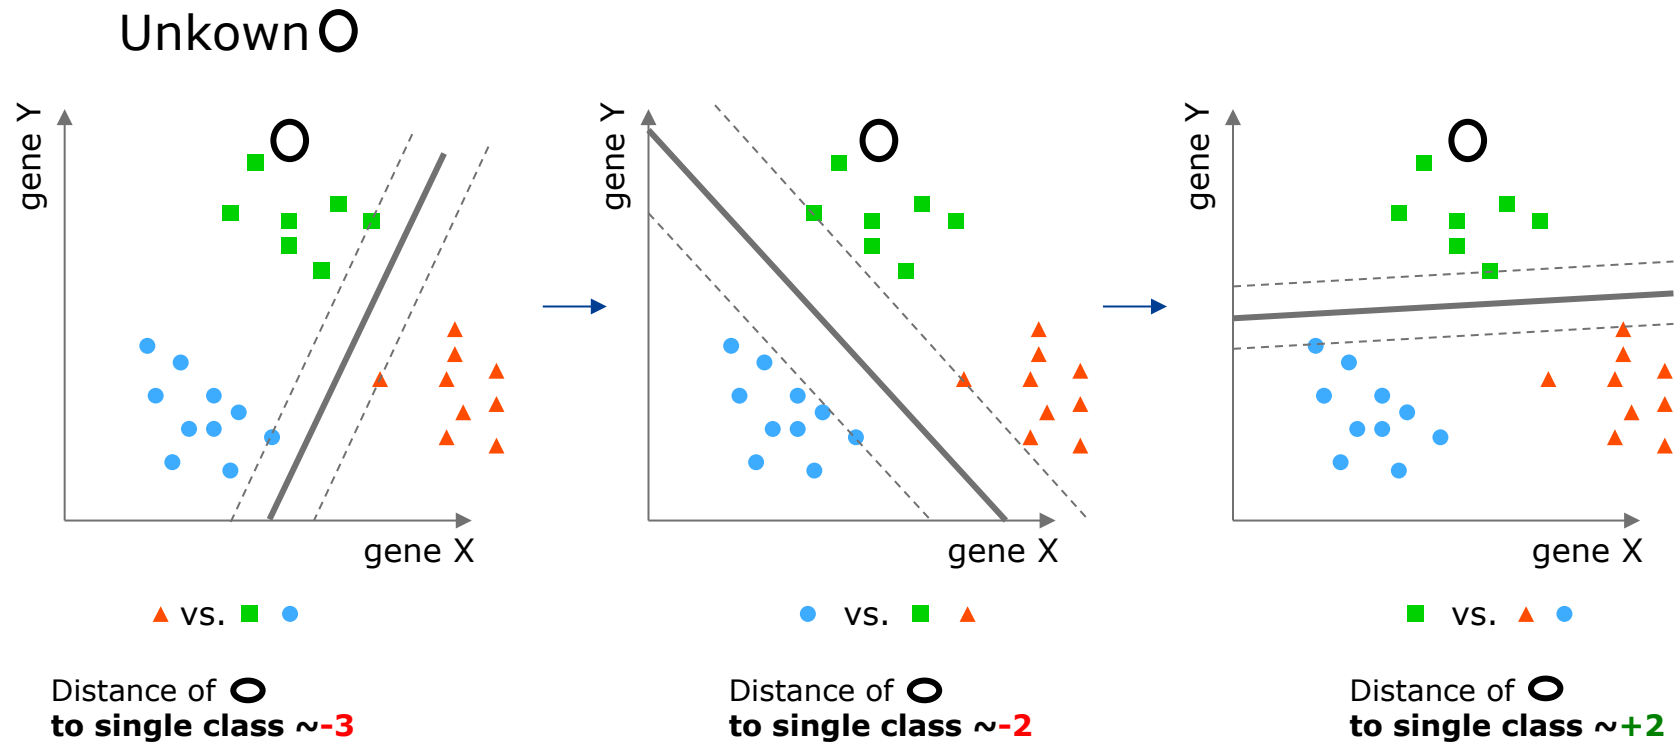
A word of warning

- It's quite fancy to write "I have used Gaussian Kernels". But always consider if you really need them!
- If number of features $>$ number of examples called ($p > n$) you probably don't need them.
- Overfitting is then the problem!
- If not it is still a good idea to try a linear kernel first!

More than 2 classes

SVM - More than 2 classes (one vs rest)

- SVM is a *binary* classifier. It can only separate two classes
- What if there are more than 2 classes?
- $N > 2$ classes N times '**one vs. rest**'




\bigcirc has the highest distance in the green case.
It will be classified as green.

One vs. all classification

```
#####  
# More than 2 classes  
# (Performs one vs all classification)  
fit = svm(Species ~ ., data=iris,  
kernel="linear", cost=10)  
res = predict(fit, iris)  
sum(res == iris$Species)
```

SVM Advanced Topics

- Custom Kernels e.g. for text
- SVM Regression
- Outlier Detection with one-class SVM (see below)



Praktikum

Bewertete Hausaufgabe

- Mitmachen an einer Data Science Challenge
- Erste Möglichkeit [Otto Produkt Klassifikation](https://www.kaggle.com/c/otto-group-product-classification-challenge) (<https://www.kaggle.com/c/otto-group-product-classification-challenge>)



- Einreichen unter:
 - http://srv-lab-t-864/submission/Otto_2016/
- Leaderboard:
 - http://srv-lab-t-864/leaderboard/Otto_2016/
- Andere Challenges von Kaggle
 - Nach Rücksprache können Sie auch an einer anderen Kaggle Challenge teilnehmen (nicht Titanic)
 - Zum Beispiel: MNIST
 - Beachten Sie, es muss ein Klassifizierungsproblem sein.
 - Username muss dann mitgeteilt werden

Bewertete Hausaufgabe

- 2er Teams OK
- Teams melden bis 9 Dezember
- Vorstellung im letzten Praktikum (20.12.2016)
 - Etwa 10-20 Minuten
- Einreichen der Lösung
- Bewertung in halben Noten
 - Performance
 - Vortrag
 - Folien
- Note zählt nur zur Verbesserung!

Code für LSG

```
X_Train = read.table("train_otto.csv", sep=';', header = TRUE, stringsAsFactors = FALSE)
X_Test = read.table("test_otto.csv", sep=';', header = TRUE, stringsAsFactors = FALSE)

# LDA
library(MASS)
fit = lda(target ~ ., data = X_Train)
res = predict(fit, X_Test)
df = data.frame(key=X_Test$id, value=res$class)
write.table(x=df, file = 'predictions_lda.csv', sep=';', row.names = FALSE)
```